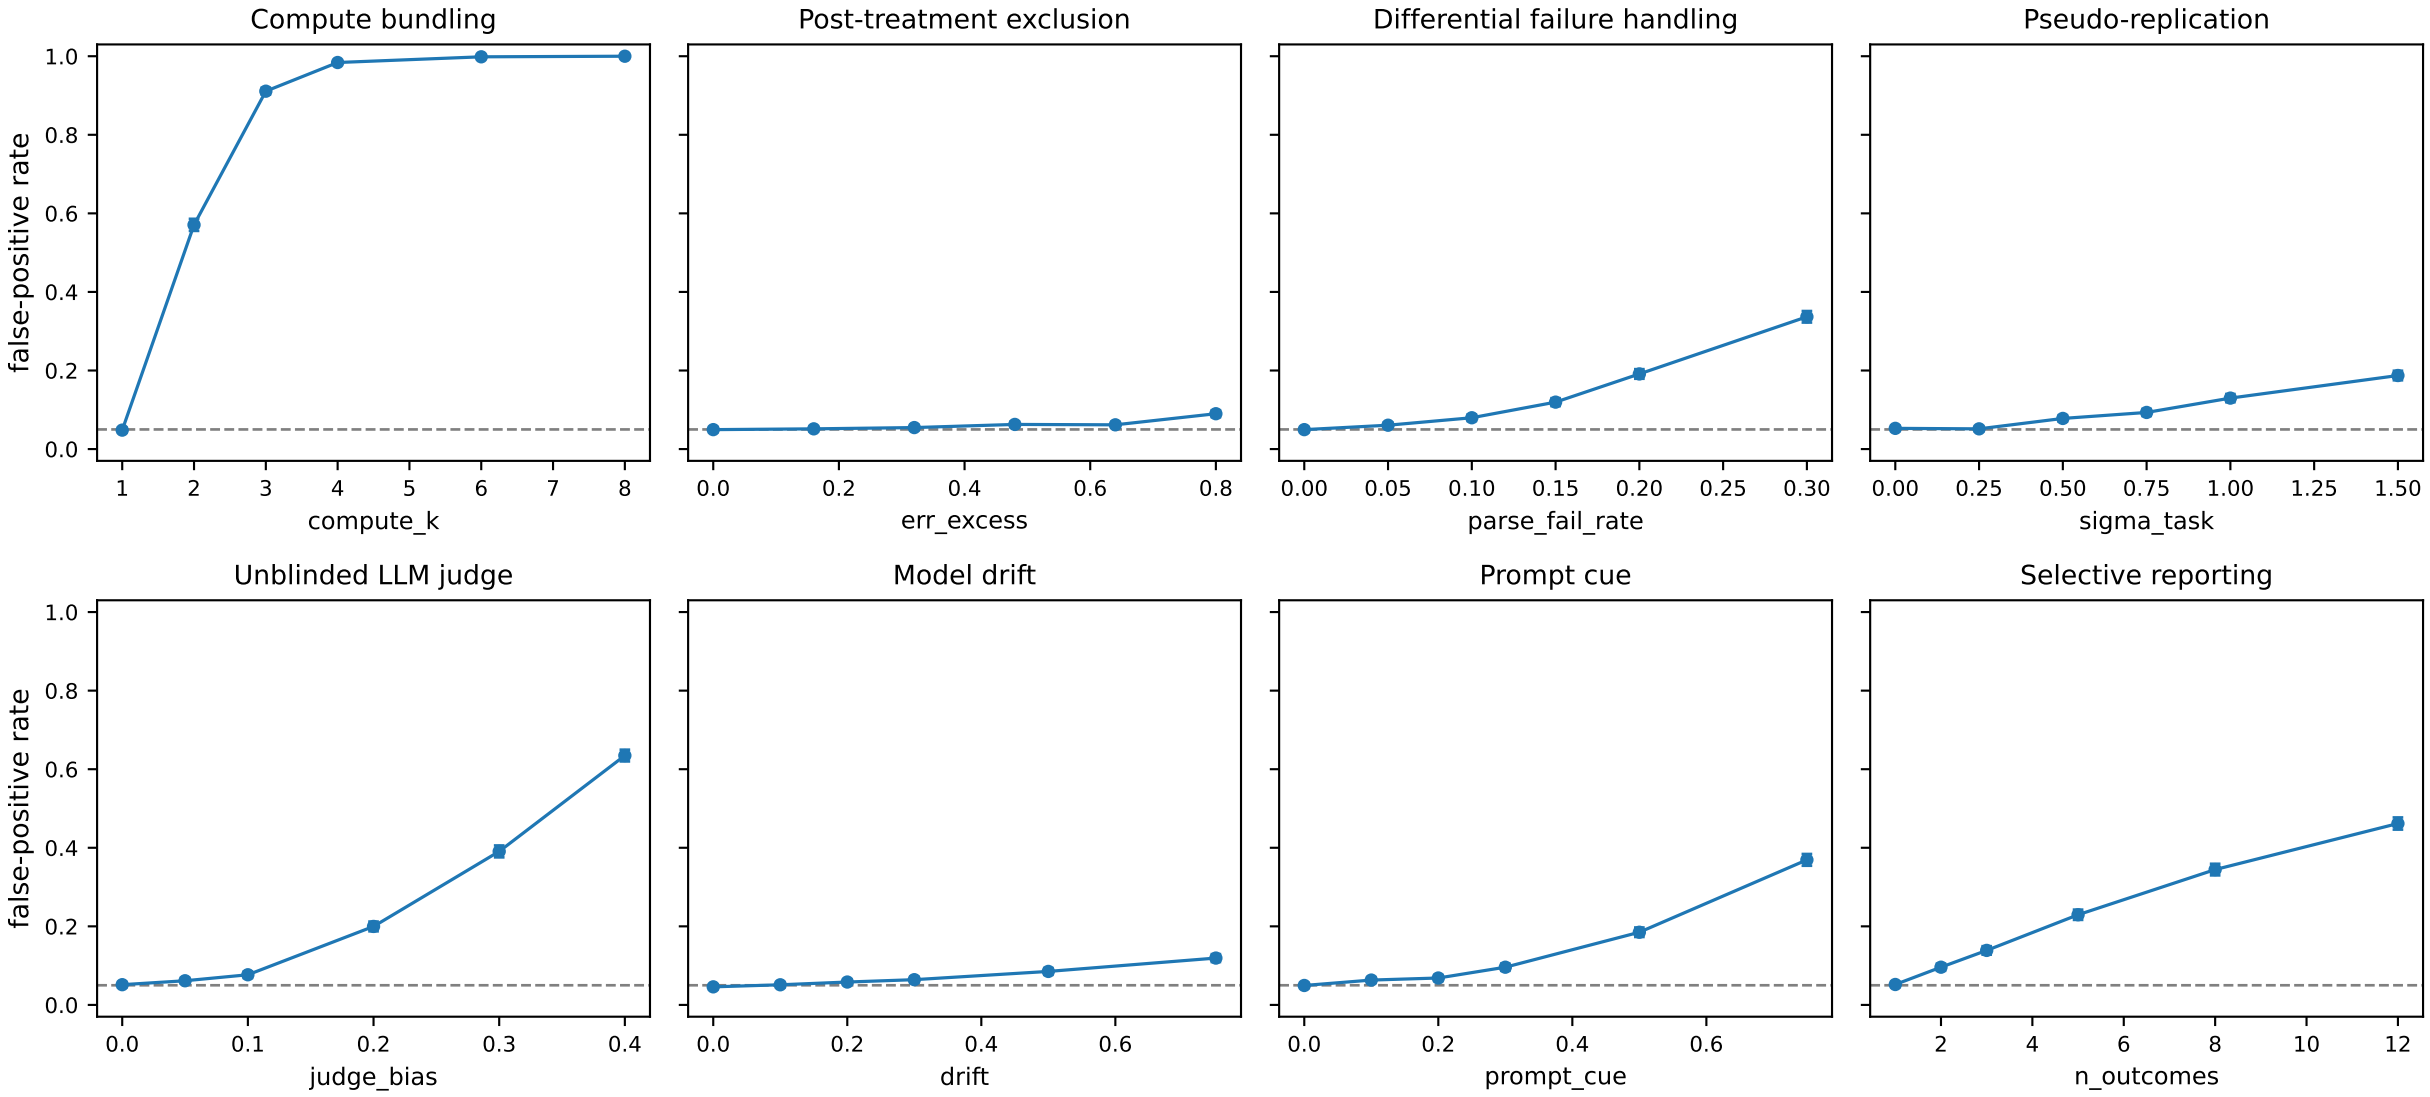
False-positive rate vs hazard strength (true effect = 0; dashed line = nominal 5%)
Compute bundling
Post-treatment exclusion
Differential failure handling
Pseudo-replication
Unblinded LLM judge
Model drift
Prompt cue
Selective reporting
false-positive rate
compute_k
err_excess
parse_fail_rate
sigma_task
judge_bias
drift
prompt_cue
n_outcomes